Handout 04: Linear Regression and Normal Equations

Linear models are amongst the most well known and often-used methods for modeling data. They are employed to study the outcomes of patients in clinical trials, the price of financial instruments, the lifetimes of fruit flies, and many other responses from a wide range of fields. Why are linear models so popular? One important attribute is that linear models provide a concrete interpretation for all of their parameters. Take the two variable model for predicting housing sale prices as a function of total area (in square feet or square meters) and the number of bedrooms,

$$\text{price}_i = \beta_0 + \beta_1 \cdot \text{area}_i + \beta_2 \cdot \text{bedrooms}_i + \epsilon_i. \tag{4.1}$$

The parameters in this model tell us how much the response, price, changes when one of the predictor variables changes with the other variable held fixed. Mathematically, we can describe this precisely using partial derivatives

$$\beta_1 = \frac{\partial \, \text{price}}{\partial \, \text{area}}, \tag{4.2}$$

$$\beta_2 = \frac{\partial \, \text{price}}{\partial \, \text{bedrooms}}. \tag{4.3}$$

The model separates the effect of the total size of a house and the total number of bedrooms. This information is useful to real estate agents, homeowners, construction companies, and economists. Linear models also allow for the interpretation of categorical predictors through the use of indicator variables. If our housing price data also includes information about whether a given observation is from one of three neighborhoods, say 'uptown,' 'downtown,' and 'suburbia,' we can define variables that are one when observation i is in the given neighborhood and zero otherwise. A linear model with these variables may be written as

$$\text{price}_i = \beta_0 + \beta_1 \cdot \text{area}_i + \beta_2 \cdot \text{bedrooms}_i + \tag{4.4}$$
$$\beta_3 \cdot \text{downtown}_i + \beta_4 \cdot \text{uptown}_i + \epsilon_i.$$

The parameter β_3 can still be viewed as a partial derivative, here representing the difference in the expected price between a house in suburbia and a house in the downtown neighborhood, if both are the same size and have the same number of bedrooms.

The relatively simple form of linear models allows for a great deal of variation in the model assumptions. The x_i's can be treated as fixed values, a *fixed design*, or they may be considered to be random variables themselves, as in a *random design* model. In biological applications the analysis usually depends on strict independence between the errors. In time series data, as commonly seen in finance or macroeconomics, the ϵ_i are often serially correlated with one another. Linear models such as

the autoregressive integrated moving average (ARIMA) model and the autoregressive conditional heteroskedasticity (ARCH) model are used to model time series data with serial correlation structures. Longitudinal medical studies, where data is collected on multiple instances from the same cohort of patients over a period of time, may assume that the errors for observations from the same subject correlate differently than errors between different patients. Fixed, random, and mixed effects models — core statistical methods within certain sub-disciplines in the sciences and social sciences — are forms of linear models adapted to handle applications such as resampled data.

Linear models also benefit from a strong theoretical background. The standard estimators, which we will explore today, can be described in terms of weighted sums of the original data. Under weak assumptions, we can then draw on the central limit theorem and large sample theory to construct asymptotically valid confidence intervals and hypothesis testing frameworks. Importantly, most of this theory can be extended to the various extensions and complex assumptions often used in practice. Also, these theoretical tools are useful even when the primary task is one of prediction. Hypothesis tests aid in the process of deciding whether to add or delete a certain variable from a model. Confidence intervals, when combined with an estimate of the noise variance, are extensible to prediction intervals. These provide a range of likely values for newly observed data points, in addition to a singular 'best' value. We will see several ways in which these estimates are useful in practice when building predictive models.

The standard estimators for parameters in linear models can be calculated using relatively straightforward computational approaches. For this reason, linear models are often used in applications even when many of the aforementioned benefits do not directly apply. Notice that a linear model must be linear only relative to the β terms. If we have pairs of data (x_i, y_i) but believe that there is a non-linear relationship between x and y, we could build the model

$$y_i = \beta_0 + \beta_1 \cdot x_i + \beta_2 \cdot x_i^2 + \cdots + \beta_p \cdot x_i^p + \epsilon_i. \qquad (4.5)$$

Here it is difficult to discern a conceptual interpretation of each of the β_j terms. As a result, it is also hard to make use of confidence intervals and hypothesis tests concerning them. However, the linear model framework is incredibly useful as it provides a computationally tractable way of estimating an arbitrarily complex relationship, by setting p as large as possible, between our two variables. Of course, the size of the dataset will limit the ultimate complexity of the model, but this is true regardless of the particular approach taken. We will expand at length on this variable expansion in the coming months.

Ordinary least squares

Many of the advantages of linear models concern the beneficial properties of the standard estimators used to compute the unknown parameters β_j from observed

data. As a next step we would like to explore the definition of these estimators. To this aim, it will be useful to provide a compact matrix-based description of a linear model. Throughout my notes, unless otherwise noted, we use a notation where n is the sample size, p is the number of variables, i is an index over the samples, and j is the index over the variables. With this notation a complete general description of a linear model can be given by

$$\widehat{y}_i = \beta_1 \cdot x_{i,1} + \cdots + \beta_p \cdot x_{i,p}, \quad \forall\, i = 1, \ldots, n. \tag{4.6}$$

Or simply

$$\widehat{y}_i = \sum_j \beta_j \cdot x_{i,j}, \quad \forall\, i = 1, \ldots, n. \tag{4.7}$$

Notice that we do not need to include an explicit intercept term β_0. If one is required this can be included by setting $x_{i,1}$ equal to one for every single observation i. Using matrix notation, we can write the linear model equation simultaneously for all observations as

$$\begin{pmatrix} \widehat{y}_1 \\ \widehat{y}_2 \\ \vdots \\ \widehat{y}_n \end{pmatrix} = \begin{pmatrix} x_{1,1} & x_{2,1} & \cdots & x_{p,1} \\ x_{1,2} & \ddots & & x_{p,2} \\ \vdots & & \ddots & \vdots \\ x_{1,n} & x_{2,n} & \cdots & x_{p,n} \end{pmatrix} \begin{pmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{pmatrix} \tag{4.8}$$

which can be compactly written in terms of a vector $\widehat{y}$ of the responses, a matrix X of the predictor variables, and a vector β of the unknown parameters

$$\widehat{y} = X\beta. \tag{4.9}$$

Beyond compactness, this notation is also useful as many of the computational properties of linear models can be reduced to linear algebraic properties of the matrix X.

Now, holding to the framework of predictive modelling, we want to know how to find a good set of values for β. To do this, we first need to define a loss function $\mathcal{L}$ that describes how well a prediction is able to predict values of y. Here we use mean squared error:

$$\mathcal{L}(\widehat{y}, y) = \sum_i (y_i - \widehat{y}_i)^2. \tag{4.10}$$

Notice that we can re-write this in matrix form following:

$$\mathcal{L} = ||y - \widehat{y}||_2^2. \tag{4.11}$$

From which we can input our formula for $\widehat{y}$

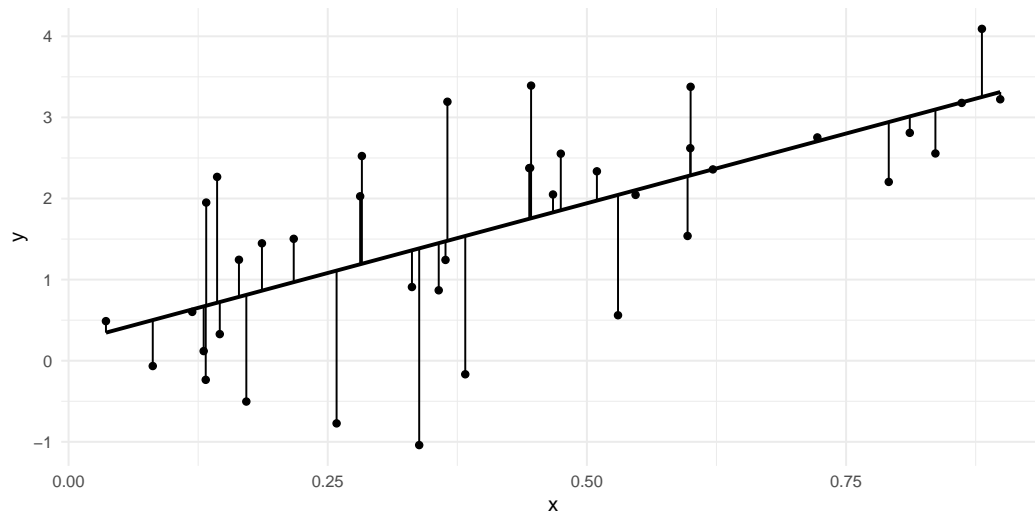$$\mathcal{L} = ||y - X\beta||_2^2. \tag{4.12}$$

Figure 1: Visualization of residuals from the linear model $y = \beta_0 + \beta_1 x$.

It is useful to also define the residual vector r, the thing to minimized in the loss function:

$$r = y - X\beta. \tag{4.13}$$

From here, the next steps are:

1. Expand the definition of the loss function.

2. Take the gradient of $\mathcal{L}$ using the matrix formulae from the last notes.

3. Set the gradient equal to zero and find the optimal value of β from the training data.

You will work through these steps today in the lab questions.

LAB QUESTIONS

1. Write Equation 4.12 as an inner product. Expand and distribute the terms so that you have a loss function written a linear combination of matrix products.

2. Convince yourself that the matrix $X^t X$ is equal to its own transpose.

3. Use the gradient rules we had last time to compute the gradient of the loss function for linear regression.

4. Set the gradient equal to zero. The result is known as the *normal equations*. Isolate β on one side using the matrix inverse.

5. Now, with a known quantity for β, write down an equation for $\widehat{y}$. This should take the form:

$$\widehat{y} = Hy \tag{4.14}$$

   For some matrix H. The matrix here is called the "hat" matrix because it puts a hat on the quantities of y. Show that $H^2 = HH = H$.

6. Assume that the "true" value of y is given by:

$$y = Xb + \epsilon \tag{4.15}$$

   For some vector of random errors ϵ. Plug this into your equation for β and show that β can be written as b plus another term that should be small if the noise terms are small.

7. Show that the residuals $(y - X\beta)$ are perpendicular to the fitted values $\widehat{y}$. That is, show that the dot product between the two is zero. (Note: you can use the fact that the transpose of $(X^t X)^{-1}$ is equal to itself).

8. Continue by following the notes in the `class04.Rmd`